

A Technical Blueprint for "Culinary Vision"

Executive Summary

This report presents a comprehensive technical blueprint for "Culinary Vision," a food recognition service designed to identify and analyze food items from user-uploaded images. The proposed solution is architected as a multi-component system comprising a sophisticated computer vision pipeline, a hybrid on-device and cloud deployment strategy, and a robust MLOps framework for continuous improvement. The project is highly feasible, but its success depends on a nuanced approach to address the unique challenges of multi-food plates and culturally specific cuisines. The Minimum Viable Product (MVP) is defined as a cross-platform mobile application that leverages a streamlined version of this architecture. Its primary objective is not only to provide a functional service but also to act as a data collection engine. The material for a future thesis will focus on the empirical validation of the continuous training loop and the performance of specialized models on a bespoke dataset of Indian and Asian cuisines. Key recommendations include using the YOLOv7 model for primary object detection in the MVP due to its speed, a lightweight MobileNetV2 backbone for on-device classification, and the more precise Mask R-CNN for the academic component of instance segmentation. The application will be developed using Flutter for cross-platform compatibility, and its long-term viability will be ensured by a data-driven MLOps pipeline that autonomously improves the model over time.

1. The Problem Space: The Intricate Challenge of Food Recognition

The domain of food image analysis is fraught with technical complexities that distinguish it from standard object recognition tasks. Unlike manufactured products with uniform shape, size, and appearance, food items exhibit significant variability, making them a particularly challenging subject for computer vision models.

1.1. Inherent Complexities in Food Image Analysis

A primary challenge lies in the diverse and fluid nature of food itself. The appearance of a single dish can vary dramatically due to differences in cooking methods, presentation styles, and even the way ingredients are cut. This creates a high degree of intra-class variation, where the same food item can look completely different in two separate images, making universal classification difficult.¹ Furthermore, food is often presented on plates with multiple, overlapping items, leading to the problem of

occlusion. A model must not only identify a food item but also accurately segment it from its background and surrounding foods, a task that simple object detection models often struggle with.² The problem is further compounded by the inherent similarity between different ingredients and the absence of a standardized definition for food items in many public datasets, which can lead to misclassification and ambiguity.²

1.2. Specifics of Asian and Indian Cuisine Recognition

The challenges outlined above are magnified when applied to Asian and Indian cuisine. India alone possesses a vast diversity of regional cuisines, where a single dish name like "biryani" can refer to many different preparations across various states, each with a unique visual signature.³ Such fine-grained distinctions necessitate a highly robust model capable of differentiating between visually similar dishes.

A thorough analysis of the problem reveals a critical finding: the project's success is not primarily contingent on the choice of the ideal model architecture, but rather on the ability to build and maintain a high-quality, culturally specific dataset. The inherent variability and fine-grained differences within these cuisines mean that pre-existing, general-purpose datasets are insufficient. Therefore, the core of this project—from the MVP to the final thesis—must revolve around a methodology for acquiring, curating, and continuously improving a custom dataset. The MVP, in this light, functions as more than just a product; it is a vital tool for data collection, providing the raw material for the long-term, self-improving system.

2. The Computer Vision Pipeline: From Pixels to Nutrition

The proposed "Culinary Vision" system employs a multi-stage deep learning architecture to transform a single image into actionable nutritional information. This pipeline is designed to sequentially address the complexities of multi-food plates and fine-grained classification.

2.1. The Multi-Stage Pipeline

The core technical workflow is divided into five sequential phases:

1. **Image Acquisition & Preprocessing:** The user-uploaded image is first processed to be compatible with the deep learning model. This includes standardization techniques such as resizing and normalization to ensure consistent input quality and improve model robustness.⁵
2. **Object Detection & Instance Segmentation:** This is the foundational step for handling multi-food plates. The model locates and separates each food item by drawing bounding boxes or, more precisely, creating pixel-wise masks around them.
3. **Fine-Grained Classification:** Each isolated food item is then passed to a separate classifier to identify its specific category (e.g., "plain rice," "dal makhani," or "chicken curry").
4. **Portion Size Estimation:** The volume or mass of each segmented food item is estimated. This step is a prerequisite for accurate calorie and macronutrient calculations.⁶
5. **Nutrition Lookup:** Finally, the classified food items and their estimated portion sizes are mapped to a comprehensive nutritional database, which provides the final output to the user.⁷

2.2. Object Detection and Instance Segmentation for Multi-Food Plates

The presence of multiple, overlapping food items on a single plate is a fundamental challenge that requires a more advanced approach than simple image classification. The solution lies in employing a model that can perform **object detection** or, for greater precision, **instance segmentation**.

A comparative analysis of state-of-the-art models highlights a clear trade-off between speed and accuracy. The **YOLO (You Only Look Once) series** is a prime candidate for the MVP due to its exceptional speed. As a "single-shot" detector, YOLO performs both classification and localization simultaneously, making it ideal for real-time inference on mobile devices.⁸ A recent

study found that YOLOv7, in particular, outperformed other algorithms in both speed and detection rate on embedded platforms.⁹

For the academic component of the project, where the objective is to achieve "peak research" and superior accuracy, **Mask R-CNN** is the more suitable choice. An extension of the Faster R-CNN architecture, Mask R-CNN is the gold standard for instance segmentation. It generates both a bounding box and a high-resolution segmentation mask for each object instance.¹⁰ A paper specifically notes its use for "solid-volume food instance segmentation and calorie estimation," confirming that it outperforms other methods by leveraging the precision of its mask information.¹¹ This dual-model approach provides a clear path for the project: the MVP can use the fast YOLOv7 for immediate functionality, while the thesis can explore the more complex and academically valuable task of instance segmentation with Mask R-CNN.

2.3. Fine-Grained Food Classification

Training a deep learning model from scratch is both computationally expensive and data-intensive. Therefore, the most effective strategy for the classification stage is to use **transfer learning**. This technique involves fine-tuning a pre-trained model on a smaller, specific food dataset, leveraging the high-level visual features it has already learned from a massive dataset like ImageNet. This approach dramatically reduces the data and computational resources required for training.⁵

For the on-device component of the MVP, a lightweight model is essential to ensure fast inference and minimal battery consumption. The **MobileNetV2** architecture is specifically designed for this purpose, with a simple structure and fewer parameters that still achieve "commendable accuracy in a short time".¹³ It has been shown to "significantly outperform the traditional CNN model" for real-time dietary monitoring on mobile devices.¹⁴

For the thesis, a more powerful model can be explored. The **EfficientNet** family of models, particularly EfficientNetB7, has demonstrated "remarkable performance" in a study, achieving high accuracy in a classification task.¹⁵ This model is deeper than others, and while it may have a longer training time, its superior performance offers a clear avenue for research into the boundaries of accuracy.¹⁵

Table 2.2.1: Comparative Analysis of Core Computer Vision Models

Model	Primary Use Case	Key Advantages	Key Disadvantages
YOLOv7	Real-time Object Detection (MVP)	Unparalleled speed; optimized for embedded platforms and real-time inference. ⁹	May struggle with small or closely-packed objects; less precise than segmentation models. ⁸
Mask R-CNN	Instance Segmentation (Thesis)	Generates pixel-accurate masks, crucial for precise calorie estimation; gold standard for segmentation tasks. ¹¹	High computational cost and latency; requires more powerful hardware for inference. ¹⁰
MobileNetV2	On-Device Classification (MVP)	Lightweight with a simpler structure; low computational and memory footprint; ideal for mobile deployment. ¹³	Less accurate than larger models; may not perform well on complex, fine-grained classification tasks. ¹⁴
EfficientNetB7	High-Accuracy Classification (Thesis)	Achieves "remarkable performance"; built on the foundation of compound scaling for efficiency. ¹⁵	Significantly longer training time; requires high-end hardware and resources. ¹⁵

3. The Minimum Viable Product (MVP) Architecture

The MVP's architecture is meticulously designed to meet the user's requirements for a

multi-platform mobile application that is robust, scalable, and cost-effective.

3.1. Hybrid On-Device and Cloud AI Strategy

A purely on-device model would be limited in its scope, while a purely cloud-based model would suffer from latency, high costs, and privacy concerns.¹⁶ Therefore, the optimal solution for the MVP is a

hybrid approach. This strategy leverages the strengths of both environments to deliver an optimal user experience while managing resources effectively.

The proposed MVP workflow begins with on-device AI. A lightweight model like MobileNetV2 handles basic classification and object detection for common food items. This provides an ultra-low latency, real-time user experience that works offline and keeps sensitive data on the device, addressing key privacy and speed requirements.¹⁶ For more complex analyses, or when the on-device model's prediction confidence is low, the system seamlessly transitions to a cloud-based fallback. The image is sent to a more powerful, larger model (e.g., Mask R-CNN) for high-precision inference and nutritional lookup.¹⁶ This tiered approach is a strategic decision for scalability and financial viability; by offloading simple, high-volume tasks to the user's device, the MVP significantly reduces the number of expensive cloud API calls.¹⁸

3.2. Cross-Platform Frameworks: A Technical Deep Dive

To meet the requirement of a single codebase for a multi-platform application (Android, Mac, Windows), a cross-platform framework is a necessity. Both Flutter and React Native are leading candidates.¹⁹

Flutter offers a strong advantage for this project. Written in Dart, it compiles directly to native ARM code, delivering "near-native performance" and a high-performance rendering engine with GPU acceleration. This is particularly beneficial for the computationally intensive tasks and graphics-rich visualizations required by an AI-powered app. Furthermore, Flutter's tight integration with Google's ecosystem and services like Firebase ML APIs streamlines development.¹⁹

React Native, in contrast, benefits from the larger and more mature JavaScript ecosystem. While it may not offer the same raw performance as Flutter for heavy computations, it

provides the flexibility to offload these tasks to optimized native modules.¹⁹

Given the reliance on TensorFlow and its associated ecosystem (MobileNetV2, etc.), **Flutter's** seamless performance and first-party support for AI integrations make it a superior choice for this project.

Table 3.2: Cross-Platform Framework Comparison for an AI App

Framework	Core Language	AI Performance	AI Integration	Pros for "Culinary Vision"
Flutter	Dart	Superior, due to native ARM compilation and GPU acceleration. ¹⁹	Tight integration with Google's Firebase ML APIs and a growing ecosystem. ¹⁹	High performance for real-time visualization and computation; streamlined development with hot-reload. ¹⁹
React Native	JavaScript	Requires native modules for heavy computations; generally lower performance. ¹⁹	Access to a vast JavaScript library ecosystem, including TensorFlow.js. ¹⁹	Large community and mature ecosystem; flexible integration of native libraries for specific tasks. ¹⁹

3.3. Nutritional Data Integration: A Review of Reliable APIs

A critical component of the MVP is the nutritional lookup, which requires a reliable and redundant API. A major constraint is the scarcity of robust databases with comprehensive data for Indian and Asian cuisines. While the USDA FoodData Central is a highly reliable source for US-based foods, its coverage of regional Indian and Asian dishes is likely limited.²¹

A review of alternatives highlights several options:

- **Ind Nutrient API:** A niche, authorization-free API that specifically provides nutritional information for Indian dishes.²³ While it is a good starting point for the MVP, its long-term reliability and redundancy are unknown.
- **Edamam:** A leading enterprise-grade provider with a massive database of over 900,000 foods and 2.3 million recipes. It is highly reliable and offers extensive nutritional data for a wide range of foods, including support for multiple languages, but it operates on a subscription model that can be costly at scale.²⁴
- **LogMeal:** An API that integrates its own AI food recognition and nutritional analysis. It could function as a full-stack solution, but its specific coverage of Indian and Asian dishes requires further investigation.²⁷

For the MVP, a tiered approach is recommended. The project can initially rely on the Ind Nutrient API to validate the core service for a limited set of Indian foods. However, to build a production-ready product and support a formal thesis, the project must either build a custom database or plan to license data from a robust provider like Edamam.

Table 3.3: Nutritional API Comparison for Indian/Asian Cuisine

API Name	Data Coverage	Cost	Reliability & Redundancy	Verdict for the MVP
Ind Nutrient API	Indian dishes. ²³	Free / Authorization-free. ²³	Unknown; not guaranteed for long-term viability. ²³	Good for an initial MVP to validate the idea, but not a long-term solution.
Edamam	Global; includes 900,000	Tiered plans, from a free trial to	High; used by major companies like	The best choice for a production-rea

	foods and 2.3 million recipes. ²⁴	enterprise-level. ²⁶	Amazon and Microsoft. ²⁵	dy product, but requires a budget.
USDA FoodData Central	Primarily U.S. foods. ²¹	Free. ²¹	High; maintained by the U.S. government. ²¹	Limited in scope for Indian/Asian foods; not suitable for a project focused on this domain. ²²

4. The Continuous Improvement Pipeline (MLOps)

The project's longevity and accuracy are dependent on a sustainable, self-improving system. This requires a robust MLOps pipeline that automates the process of data acquisition, model retraining, and deployment.

4.1. Data Acquisition and Curation Strategy

Research indicates that the lack of a large, high-quality, and annotated dataset is a significant bottleneck in food recognition.² Therefore, a critical component of this project is the creation of a custom dataset. This process begins by leveraging existing public datasets to bootstrap the initial model.

- **Indian Food Images Dataset (Kaggle):** A solid starting point for classification, with 4000 images across 80 classes.²⁹
- **UEC-FoodPix Complete:** An excellent resource for training the instance segmentation model, containing 10,000 images with manually annotated segmentation masks.³⁰
- **COCO Dataset:** A large, general-purpose object detection dataset that includes some food categories, useful for initial training of the object detection model.³¹

For custom annotation, an open-source platform like **Label Studio** is highly recommended. It supports all necessary annotation types, including bounding boxes and polygons for segmentation. Its "ML-assisted labeling" feature can significantly speed up the annotation

process by using initial model predictions to suggest labels for human review.³² The data collection method will combine publicly scraped images with crowd-sourced images captured directly from the mobile application, which will serve as a live sensor for new data.³³

Table 4.1.1: Summary of Relevant Public Datasets for Indian/Asian Food

Dataset Name	Size & Scope	Annotation Type	Project Relevance
Indian Food Images Dataset	4000 images, 80 classes of Indian food. ²⁹	Image Classification	Provides a strong foundation for the fine-grained classification component of the MVP.
UEC-FoodPix Complete	10,000 images, 103 food categories. ³⁰	Instance Segmentation Masks	Essential for training the Mask R-CNN model for the thesis-level research.
MS COCO	>330K images, 80 categories. ³¹	Bounding boxes, Segmentation masks	Useful for initial training of object detection models due to its size and diversity. ³¹

4.2. The Active and Continual Learning Loop

The user's concept of using new images to train the model is a form of **continual learning**, a critical MLOps practice that allows a model to learn from a constant stream of new data without "catastrophic forgetting".³⁴ This closed-loop system ensures the model remains accurate and adapts to "data drift," such as new dishes or preparation styles.³⁵

The proposed pipeline operates as follows:

1. **Prediction and Monitoring:** The deployed MVP's predictions are monitored in real-time. Key metrics, such as prediction confidence and user feedback, are tracked by a "Model Monitoring Service".³⁵
2. **Data-Driven Retraining:** The MLOps pipeline is automatically triggered when the monitoring service detects anomalies, such as a high volume of low-confidence predictions or poor user feedback on a specific category.³⁵
3. **Active Learning for Labeling:** Instead of manually labeling all new data, an **Active Learning** strategy is used to select the "most informative" and "uncertain" samples for human annotation.³⁶ This minimizes the manual labeling burden and maximizes the model's learning efficiency by focusing on the most challenging examples.
4. **Automated Retraining & Deployment:** The newly labeled data is added to a feature store. An automated pipeline retrains the model, evaluates a new "contender model," and deploys it to the production endpoint if it meets predefined performance metrics.³⁵

This process transforms the MVP into a powerful, automated data collection engine. Its primary function is to capture user-generated images and feedback, which serves as a live source of ground truth data. This addresses the core data scarcity problem and provides the unique, proprietary dataset required for a robust, high-performing product and a compelling academic thesis.

5. Material for a Thesis: Advanced Research & Future Work

The project provides a fertile ground for a formal academic thesis, pushing beyond the capabilities of the MVP to address key research challenges and explore future directions.

5.1. Overcoming Key Technical Hurdles

While the MVP can use a simplified method for calorie estimation (e.g., standard portion sizes), the problem of **portion size estimation** from a single image is a major, unsolved research challenge in the field.⁶ The thesis can propose a more advanced solution by incorporating a known reference object (e.g., a credit card or a coin) in the image for scale calibration. This would require developing complex geometric algorithms to convert 2D image coordinates to 3D world coordinates, a task that has been the focus of prior research.³⁸ By combining this with the precise instance segmentation from Mask R-CNN, the thesis can

propose a method for a more accurate volume and calorie estimation system.

5.2. Proposed Research Questions for an Academic Thesis

The research can be formalized by addressing specific, high-level questions that advance the field of food recognition:

- **Question 1:** To what extent can an instance segmentation model (e.g., Mask R-CNN) improve the accuracy of portion size estimation from a single, monocular image of a multi-food plate?
- **Question 2:** How does a continual learning pipeline, augmented by an active learning sampling strategy, impact the long-term accuracy and efficiency of a food recognition model on a specific, culturally-defined dataset?
- **Question 3:** What is the optimal balance between a lightweight on-device model and a high-accuracy cloud-based model to optimize both user experience and long-term data collection efficiency?

5.3. A Vision for a Comprehensive Dietary Assessment System

Beyond the initial scope, the project can evolve into a comprehensive dietary assessment tool. Future research could explore more granular tasks, such as **image-based ingredient identification** from a single dish, which is a more complex problem than multi-dish segmentation.² The system could also be enhanced with

3D image reconstruction for more accurate volume estimation, which has been explored in academic literature.⁶ Finally, the model could benefit from

multi-modal learning, which combines visual data with text (e.g., a user-provided dish name) to improve classification accuracy and overcome the inherent ambiguities of a single image.³

6. A Detailed Product Description and User Experience

The final product, "Culinary Vision," is a mobile application that simplifies dietary tracking with

a seamless, intuitive user experience.

6.1. The User Flow: From Image Capture to Insight

The user experience is designed for simplicity and speed. A user captures an image of their plate with their smartphone camera. The on-device model immediately performs a preliminary scan, providing a real-time, low-latency prediction for a limited set of common dishes. If the user desires more detailed information, they can tap on a food item to initiate a deeper analysis. This triggers the cloud-based portion of the pipeline, which uses a more powerful model to perform a high-precision, multi-item breakdown and nutritional lookup. The results are displayed clearly, showing the identified food items, their estimated calories, and other relevant details. A crucial feature of the MVP is the user feedback mechanism, which allows users to correct a misidentified food or add a new one. This human-in-the-loop feedback is the essential component that powers the continuous improvement pipeline.

6.2. Post-MVP Roadmap and Monetization Paths

The project's strategy is structured into three distinct phases:

- **Phase 1: MVP Release (Data Collection):** The initial launch will focus on a free, ad-supported application. The core value proposition is the ease of dietary tracking and a growing database of Indian and Asian cuisine. The primary objective is to acquire a critical mass of users to begin collecting a proprietary dataset with user-validated labels.
- **Phase 2: Product Expansion:** As the dataset grows, the product will be expanded with advanced features, including portion size estimation based on the research developed in the thesis, detailed dietary tracking over time, and personalized meal recommendations.
- **Phase 3: Monetization:** Once the product's accuracy and feature set are well-established, a subscription service can be offered. Premium features could include detailed micronutrient analysis, personalized diet plans, and integrations with other health and fitness applications. The unique, high-quality, and proprietary dataset will serve as a key competitive advantage in the market, enabling the creation of a truly valuable and distinct dietary assessment tool.

Works cited

1. [www.mdpi.com](https://www.mdpi.com/2076-3417/15/14/7626#:~:text=It%20is%20difficult%20to%20extract,and%20differences%20within%20food%20images.), accessed on August 30, 2025, <https://www.mdpi.com/2076-3417/15/14/7626#:~:text=It%20is%20difficult%20to%20extract,and%20differences%20within%20food%20images.>

2. Recognizing Multiple Ingredients in Food Images Using a Single-Ingredient Classification Model - IGI Global, accessed on August 30, 2025, <https://www.igi-global.com/article/recognizing-multiple-ingredients-in-food-images-using-a-single-ingredient-classification-model/360782>
3. (PDF) A MACHINE LEARNING-BASED APPROACH FOR RECOGNIZING AND CLASSIFYING TRADITIONAL SOUTH ASIAN FOOD ITEMS - ResearchGate, accessed on August 30, 2025, https://www.researchgate.net/publication/394527146_A_MACHINE_LEARNING-BASED_APPROACH_FOR_RECOGNIZING_AND_CLASSIFYING_TRADITIONAL_SOUTH_ASIAN_FOOD_ITEMS
4. A deep learning model with machine vision system for recognizing type of the food during the food consumption - PMC, accessed on August 30, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12350950/>
5. Deep Learning Based Food Recognition, accessed on August 30, 2025, <https://cs229.stanford.edu/proj2016/report/YuMaoWang-Deep%20Learning%20Based%20Food%20Recognition-report.pdf>
6. A COMPARISON OF FOOD PORTION SIZE ESTIMATION USING GEOMETRIC MODELS AND DEPTH IMAGES - PMC, accessed on August 30, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC6226035/>
7. Image-based food monitoring and dietary management for patients living with diabetes: a scoping review of calorie counting applications - PubMed Central, accessed on August 30, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11983398/>
8. 9 Best Object Detection Models of 2025: Reviewed & Compared - Hitech BPO, accessed on August 30, 2025, <https://www.hitechbpo.com/blog/top-object-detection-models.php>
9. General improved SSD model for picking object recognition of multiple fruits in natural environment - ResearchGate, accessed on August 30, 2025, https://www.researchgate.net/publication/329246944_General_improved_SSD_model_for_picking_object_recognition_of_multiple_fruits_in_natural_environment
10. matterport/Mask_RCNN: Mask R-CNN for object detection and instance segmentation on Keras and TensorFlow - GitHub, accessed on August 30, 2025, https://github.com/matterport/Mask_RCNN
11. Utilizing Mask R-CNN for Solid-Volume Food Instance ... - MDPI, accessed on August 30, 2025, <https://www.mdpi.com/2076-3417/12/21/10938>
12. Indian Food Image Recognition with MobileNetV2 - IRJET, accessed on August 30, 2025, <https://www.irjet.net/archives/V8/i8/IRJET-V8I8102.pdf>
13. (PDF) Computer Vision in the Food Industry: Accurate, Real-time, and Automatic Food Recognition with Pretrained MobileNetV2 - ResearchGate, accessed on August 30, 2025, https://www.researchgate.net/publication/380730357_Computer_Vision_in_the_Food_Industry_Accurate_Real-time_and_Automatic_Food_Recognition_with_Pretrained_MobileNetV2
14. Food Image Classification with Improved MobileNet Architecture and Data Augmentation, accessed on August 30, 2025,

- <https://www.semanticscholar.org/paper/Food-Image-Classification-with-Improved-MobileNet-Phiphaphatphaisit-Surinta/c615259b3418b8ea69d3463bd79f0b061c85d1af>
15. (PDF) A deep learning model with machine vision system for recognizing type of the food during the food consumption - ResearchGate, accessed on August 30, 2025, https://www.researchgate.net/publication/394473766_A_deep_learning_model_with_machine_vision_system_for_recognizing_type_of_the_food_during_the_food_consumption
 16. On-Device AI vs Cloud AI: Unlock Peak Performance for App, accessed on August 30, 2025, <https://appbirds.co/on-device-ai-vs-cloud-ai-best-for-your-app/>
 17. Should your mobile AI run on-device or in the cloud? - YouTube, accessed on August 30, 2025, <https://www.youtube.com/shorts/2won7XwyZwQ>
 18. TensorFlow Lite vs PyTorch Mobile for On-Device Machine Learning - ProAndroidDev, accessed on August 30, 2025, <https://proandroiddev.com/tensorflow-lite-vs-pytorch-mobile-for-on-device-machine-learning-1b214d13635f>
 19. Building Cross-Platform AI-Powered Apps with Flutter & React ..., accessed on August 30, 2025, <https://200oksolutions.com/blog/building-cross-platform-ai-powered-apps-with-flutter-react-native/>
 20. Cross Platform Hybrid Mobile App Development | React Native, Flutter - Technology Rivers, accessed on August 30, 2025, <https://technologyrivers.com/services/cross-platform-hybrid-mobile-app-development/>
 21. Food Data Central API | Documentation | Postman API Network, accessed on August 30, 2025, <https://www.postman.com/api-evangelist/agricultural-research-service-ars/documentation/nex4lq6/food-data-central-api>
 22. USDA FoodData Central, accessed on August 30, 2025, <https://fdc.nal.usda.gov/>
 23. Ind Nutrient API - RapidAPI, accessed on August 30, 2025, <https://rapidapi.com/ind-nutrient-api-ind-nutrient-api-default/api/ind-nutrient-api-1>
 24. Edamam - Food Database API, Nutrition API and Recipe API, accessed on August 30, 2025, <https://www.edamam.com/>
 25. Top 8 Nutrition APIs for Meal Planning 2024 - Eat Fresh Tech, accessed on August 30, 2025, <https://www.eatfresh.tech/blog/top-8-nutrition-apis-for-meal-planning-2024>
 26. Nutrition Analysis API - Nutrition Data Service - Edamam, accessed on August 30, 2025, <https://developer.edamam.com/edamam-nutrition-api>
 27. Recipe API - 12 Most Popular Recipe APIs [Q3, 2024] - Suggestic, accessed on August 30, 2025, <https://blog.suggestic.com/recipe-api-ultimate-list>
 28. (PDF) Computer Vision for Food Quality Assessment: Advances and ..., accessed on August 30, 2025, https://www.researchgate.net/publication/389281944_Computer_Vision_for_Food

Quality Assessment Advances and Challenges

29. Indian Food Images Dataset - Kaggle, accessed on August 30, 2025, <https://www.kaggle.com/datasets/iamsouravbanerjee/indian-food-images-dataset>
30. UEC-FoodPix Complete: A Large-Scale Food Image Segmentation Dataset - ResearchGate, accessed on August 30, 2025, https://www.researchgate.net/publication/349475604_UEC-FoodPix_Complete_A_Large-Scale_Food_Image_Segmentation_Dataset
31. COCO dataset, accessed on August 30, 2025, <https://cocodataset.org/>
32. Label Studio: Open Source Data Labeling, accessed on August 30, 2025, <https://labelstud.io/>
33. The Good, The Bad and The Ugly: An Open Image Dataset for Automated Sorting of Good, Bad, and Imperfect Produce Using AI and Robotics - MDPI, accessed on August 30, 2025, <https://www.mdpi.com/2071-1050/16/15/6411>
34. How to use continual learning to your machine learning models ..., accessed on August 30, 2025, <https://cnvrg.io/how-to-use-continual-learning-to-your-machine-learning-models/>
35. A Guide to Continuous Training of Machine Learning Models in ..., accessed on August 30, 2025, <https://www.omdena.com/blog/continuous-training-machine-learning-models>
36. [2505.06825] Active Learning for Multi-class Image Classification - arXiv, accessed on August 30, 2025, <https://www.arxiv.org/abs/2505.06825>
37. Unlabeled data selection for active learning in image classification - PMC - PubMed Central, accessed on August 30, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10764761/>
38. Determination of Food Portion Size by Image Processing - ResearchGate, accessed on August 30, 2025, https://www.researchgate.net/publication/224338657_Determination_of_Food_Portion_Size_by_Image_Processing
39. DeepFood: Automatic Multi-Class Classification of Food Ingredients Using Deep Learning, accessed on August 30, 2025, <https://web.umkc.edu/chen/PDF/CIC2017.pdf>
40. International Journal of Engineering Technology Research ... - IJETRM, accessed on August 30, 2025, <https://ijetrm.com/issues/files/Apr-2025-01-1743511350-AP01.pdf>